

RASLEN GUESMI

AI Engineer | LLM Fine-Tuning & RL | MLOps

Hammamet, Tunisia (open to EU relocation on contract) | guesmiraslen2@gmail.com | +216 40 309 961 | linkedin/raslen-guesmi | github/raslen10 | Portfolio

SUMMARY

AI/ML engineer with hands-on LLM fine-tuning experience (supervised fine-tuning of GPT-4-mini for a production use case) and applied reinforcement learning work: Q-Learning/SARSA policy experiments and live multi-agent RL trading swarms orchestrated end-to-end and tracked in MLflow. Full MLOps practice — training-data pipelines, containerised deployment, CI/CD, cloud GPU workflows on AWS/GCP — applied across fintech, health-tech, marketing, and energy domains. Comfortable owning systems end to end, from data curation through evaluation, deployment, and product analytics.

TECHNICAL SKILLS

Languages: Python, SQL

ML & Deep Learning: PyTorch, TensorFlow, Keras, Scikit-learn, LSTM, CNN, Reinforcement Learning (Q-Learning, SARSA, multi-agent RL)

LLM & Fine-Tuning: Supervised fine-tuning (OpenAI GPT-4-mini), RAG, LangChain, LangGraph, multi-agent orchestration, prompt/agent evaluation

Data Engineering: Training-data curation & feature pipelines (Airflow, PySpark, dbt), PostgreSQL, pgvector, Supabase, cleaning & deduplication at pipeline stage

Cloud & MLOps: AWS, GCP, Docker, GitHub Actions CI/CD, MLflow, Coolify, Hostinger VPS, Linux, Git, containerised GPU workflows (Azure ML/AI Foundry: ramping fast)

Product Analytics: Amplitude, Meta Pixel & Conversions API, funnel instrumentation

WORK EXPERIENCE

Senior Consultant — Evely (evely.health)

Jun 2026 – Present | Miami, USA (Remote)

- Migrated the platform from Next.js to FastHTML + FastAPI and from Vercel to a self-hosted Coolify/Hostinger VPS stack, owning the full deployment pipeline and CI/CD.
- Redesigned onboarding UI/UX and built an Amplitude analytics dashboard to track the funnel; cut drop-off from 40% to 5–10%.
- Optimised the Gemini-powered AI report feature — tuned prompts and added fallback-model handling for faster, more reliable generation.
- Restructured the PostgreSQL schema and queries to speed up dashboard and reporting performance; resolved Supabase RLS security vulnerabilities.

AI Engineer | Data Scientist — PredictiveLabs, Financial Sector

Jun 2025 – Present | Remote, London UK

- Finespresso (Jun 2025–Present): trained and evaluated Q-Learning and SARSA policies to optimise trading decisions, then built and deployed a live multi-agent reinforcement-learning swarm across four exchanges on AWS, with coordinated decision-making and risk controls, tracking every run (accuracy, R^2 , PnL) in MLflow.
- AssetHero (assethero.chat, Jan 2026–Present): built the API orchestration layer unifying four independent trading agent apps — AlpTrade, TradeSwarm, MacroHero, PolyTrade — each in its own repo, into a single callable multi-agent trading product, deployed via Docker on GCP Cloud Run.
- Tendly (2026): built scheduled data-scraping pipelines across 10+ countries, deployed as GCP Cloud Run jobs.
- POLLY (May–Aug 2026, polly.predictivelabs.ai): built a hierarchical RAG pipeline (pgvector, HNSW, chunking, embeddings, metadata) and LangChain/LangGraph agents, packaged as an MCP server, behind a compliance-aware AI marketing platform deployed via Docker on Coolify (Hostinger VPS).
- Built training-data ETL pipelines (news, market data) into a near-real-time PostgreSQL feature store, with cleaning, structuring, and feature extraction via LLM/RAG pipelines.

Freelance AI/ML Engineer — Upwork, Top Rated

Jan 2025 – Present | Remote

- Fine-tuned GPT-4-mini (supervised fine-tuning) for a digital creator client: curated and formatted training data, ran the fine-tuning job, and evaluated output against a target tone and topic set (brand deals, monetisation, mindset).
- Built structured JSON training datasets and data-prep pipelines for multiple clients' GPT fine-tuning projects.

Data Science Engineer | ML Engineer — 4InaTechnology, Energy Sector

Feb 2024 – Aug 2024 | Manouba, Tunisia

- Trained LSTM and ARIMA models from scratch on sensor and current (electrical load) data for energy consumption forecasting.
- Built the MVP end-to-end: FastAPI backend serving the model, Flutter web/mobile client, containerised with Docker and deployed on Coolify.

EDUCATION & CERTIFICATIONS

MSc Data Science, University of Bizerte, Tunisia (2021–2023) | BSc Applied Mathematics, University of Bizerte, Tunisia (2018–2021)

Intro to AI (University of Helsinki) | Intro to MLflow (DataCamp) | Power BI Job Simulation (PwC Switzerland, Forage) | Web Scraping in Python (DataCamp)

Languages: English C1 | French Fluent | Arabic Native | German A2 (targeting B1)